

Class11: RNA-Seq continued

Sam Altshuler (PID: A59010373)

2/23/2022

Transcriptomics and Analysis of RNA-Seq Data

Today we will run differential expression analysis of published data from Himes et al.

Import the countData and colData

```
counts <- read.csv("airway_scaledcounts.csv", row.names = 1)
metadata <- read.csv("airway_metadata.csv")
```

```
head(counts)
```

```
##           SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG00000000003      723       486       904       445       1170
## ENSG00000000005        0        0        0        0        0
## ENSG00000000419      467       523       616       371       582
## ENSG00000000457      347       258       364       237       318
## ENSG00000000460       96        81        73        66       118
## ENSG00000000938        0        0         1         0         2
##           SRR1039517 SRR1039520 SRR1039521
## ENSG00000000003      1097       806       604
## ENSG00000000005        0        0         0
## ENSG00000000419      781       417       509
## ENSG00000000457      447       330       324
## ENSG00000000460       94       102        74
## ENSG00000000938        0        0         0
```

There are 38694 rows, i.e. “genes” in this dataset. There are 8 columns in this dataset, i.e. experiments in the dataset.

```
metadata
```

```
##      id    dex celltype    geo_id
## 1 SRR1039508 control  N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control  N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control  N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
## 7 SRR1039520 control  N061011 GSM1275874
## 8 SRR1039521 treated   N061011 GSM1275875
```

The rows in the metadata set corresponds to the experiments being run (the columns in the counts dataset). There are 4 controls and 4 treated experiments.

The next question is does the drug do anything?

First confirm that the metadata matches the counts data.

```
#column names of counts compared to ID column of metadata  
all(colnames(counts) == metadata$id)
```

```
## [1] TRUE
```

All of the data names match up!

Gather all of the control data (extract from metadata).

```
#Store the IDs of the control experiments  
control <- metadata[metadata$dex == "control", "id"]  
# Pull the columns corresponding to the controls from the counts dataset  
ct_control <- counts[,control]
```

Gather all of the treated data. This is the same as above but for the treated columns.


```
treat <- metadata[metadata$dex == "treated", "id"]  
ct_treat <- counts[,treat]
```

Get a mean gene expression per gene for both the control and the treated. Use `apply()` or `rowMeans()`.

```
# using apply: apply(ct_control, 1, mean)  
control.mean <- rowMeans(ct_control)  
treated.mean <- rowMeans(ct_treat)
```

Compare the two experimental conditions in a plot.

```
plot(treated.mean, control.mean)
```




There are a bunch of genes with low values that overlap, making it hard to ID individual genes. The data is very skewed. The solution is to transform the data (like a log transformation) to make it more readable especially around overlapping values.

```
plot(treated.mean, control.mean, log = "yx")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 x values <= 0 omitted  
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 y values <= 0 omitted  
## from logarithmic plot
```



We need to get rid of zeros because you can't take a log of 0.

We often use log 2 transformation because it has an easier to understand output. A log 2 value of zero means that there's been no change (lies on the straight line). A value of 1 means it's doubled from treatment compared to control and -1 means it's half. This is called the fold change (how much is it doubling)

```
log2fc <- log2(treated.mean/control.mean)
```

Make a dataframe to store the results

```
ct_mean <- data.frame(control.mean, treated.mean, log2fc)
head(ct_mean)
```

##	control.mean	treated.mean	log2fc
## ENSG000000000003	900.75	658.00	-0.45303916
## ENSG000000000005	0.00	0.00	NaN
## ENSG000000000419	520.50	546.00	0.06900279
## ENSG000000000457	339.75	316.50	-0.10226805
## ENSG000000000460	97.25	78.75	-0.30441833
## ENSG000000000938	0.75	0.00	-Inf

If either the control or treated have values of zero, there won't be a valuable log2fc value. If the denominator is a zero, the answer will be NaN (not a number), and if it's in the numerator, the answer will be infinity.

Try to find and filter out the zero values.

```

# Choose all count values for both control and treated with a value of zero
# make sure to return array indices where in the dataframe the zero values are
# Save the rows that correspond to the zero values
zip <- unique(which(ct_mean[,1:2] == 0, arr.ind = TRUE)[,"row"])
# Remove the rows that correspond to zero values from the dataframe
ct_mean_2 <- ct_mean[-zip,]
head(ct_mean_2)

```

```

##               control.mean treated.mean      log2fc
## ENSG000000000003      900.75      658.00 -0.45303916
## ENSG000000000419      520.50      546.00  0.06900279
## ENSG000000000457      339.75      316.50 -0.10226805
## ENSG000000000460       97.25       78.75 -0.30441833
## ENSG000000000971     5219.00     6687.50  0.35769358
## ENSG00000001036     2327.00     1785.75 -0.38194109

```

There are 21817 genes left after removing the zero values.

There are 250 genes that have a log2fc more than +2 (upregulated).

```
sum(ct_mean_2$log2fc > 2)
```

```
## [1] 250
```

These log2fc may not actually be statistically significant. Time to use the DESeq2 package!

DESeq2

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```

##      anyDuplicated, append, as.data.frame, basename, cbind, colnames,
##      dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,
##      grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,
##      order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,
##      rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,
##      union, unique, unsplit, which.max, which.min

```

```

##
## Attaching package: 'S4Vectors'

## The following objects are masked from 'package:base':
##
##     expand.grid, I, unname

## Loading required package: IRanges

##
## Attaching package: 'IRanges'

## The following object is masked from 'package:grDevices':
##
##     windows

## Loading required package: GenomicRanges

## Loading required package: GenomeInfoDb

## Loading required package: SummarizedExperiment

## Loading required package: MatrixGenerics

## Loading required package: matrixStats

##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##     colAlls, colAnyNAs, colAnys, colAvgsPerRowSet, colCollapse,
##     colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##     colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##     colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##     colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##     colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##     colWeightedMeans, colWeightedMedians, colWeightedSds,
##     colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgsPerColSet,
##     rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
##     rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
##     rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
##     rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
##     rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
##     rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
##     rowWeightedSds, rowWeightedVars

## Loading required package: Biobase

```

```
## Welcome to Bioconductor
##
## Vignettes contain introductory material; view with
## 'browseVignettes()'. To cite Bioconductor, see
## 'citation("Biobase)", and for packages 'citation("pkgname)".

##
## Attaching package: 'Biobase'

## The following object is masked from 'package:MatrixGenerics':
##
## rowMedians

## The following objects are masked from 'package:matrixStats':
##
## anyMissing, rowMedians
```

First we need to set up the DESeq data object.

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```

```
dds
```

```
## class: DESeqDataSet
## dim: 38694 8
## metadata(1): version
## assays(1): counts
## rownames(38694): ENSG000000000003 ENSG000000000005 ... ENSG00000283120
## ENSG00000283123
## rowData names(0):
## colnames(8): SRR1039508 SRR1039509 ... SRR1039520 SRR1039521
## colData names(4): id dex celltype geo_id
```

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

```
res <- results(dds)
res
```


```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 38694 rows and 6 columns
##           baseMean log2FoldChange      lfcSE      stat      pvalue
##           <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003  747.1942    -0.3507030  0.168246  -2.084470  0.0371175
## ENSG000000000005    0.0000         NA         NA         NA         NA
## ENSG000000000419  520.1342     0.2061078  0.101059   2.039475  0.0414026
## ENSG000000000457  322.6648     0.0245269  0.145145   0.168982  0.8658106
## ENSG000000000460   87.6826    -0.1471420  0.257007  -0.572521  0.5669691
## ...           ...           ...           ...           ...
## ENSG00000283115    0.000000         NA         NA         NA         NA
## ENSG00000283116    0.000000         NA         NA         NA         NA
## ENSG00000283119    0.000000         NA         NA         NA         NA
## ENSG00000283120    0.974916    -0.668258   1.69456  -0.394354  0.693319
## ENSG00000283123    0.000000         NA         NA         NA         NA
##           padj
##           <numeric>
## ENSG000000000003  0.163035
## ENSG000000000005         NA
## ENSG000000000419  0.176032
## ENSG000000000457  0.961694
## ENSG000000000460  0.815849
## ...           ...
## ENSG00000283115         NA
## ENSG00000283116         NA
## ENSG00000283119         NA
## ENSG00000283120         NA
## ENSG00000283123         NA
```

padj is the adjusted p-value for multiple testing.

A main results figure

A common main results figure is a volcano plot. This is a plot of the log2 fold change on the x axis v the p-value (or padj) on the y-axis.

```
plot(res$log2FoldChange, -log(res$padj))
```





As it goes up the y-axis, the smaller the p-value, the less likely the fold change is due to random chance (i.e. false positiveness). The plot function should be `plot(foldchange, -log(p-value))`.

```
library(ggplot2)
# make the same volcano plot as above but color by if the p value is less than 0.05 AND
# the log2 fold change is greater than 2 or less than -2 (absolute value is greater than 2)
ggplot(as.data.frame(res))+
  aes(x = log2FoldChange, y = -log(padj), color = padj < 0.05 & abs(log2FoldChange) > 2)+
  geom_point()+
  xlab("Log2 Fold Change") +
  ylab("-Log(Adjusted p-Value)") +
  labs(title = "Differential Gene Expression", caption = "Data from Himes et al.") +
  theme(legend.position = "none")
```

```
## Warning: Removed 23549 rows containing missing values (geom_point).
```

Differential Gene Expression



Data from Himes et al.